

Mosh

An Interactive Remote Shell for Mobile Clients

Keith Winstein (with Hari Balakrishnan)

keithw@mit.edu

May 31, 2012

Secure Shell, 1995

- ▶ Connects local terminal to remote terminal.
- ▶ Conveys over TCP:
 - ▶ user keystrokes → server
 - ▶ octet stream (coded screen updates) → client terminal
- ▶ Connection endpoints dictated by IP:port on both sides

Post-1995 problems with SSH

- ▶ Can't roam:
 - ▶ ... across Wi-Fi networks.
 - ▶ ... from Wi-Fi to cell or vice versa.
- ▶ TCP times out if data unacknowledged after n minutes.
 - ▶ Session dies if laptop goes to sleep.
- ▶ TCP responds poorly to packet loss.

More problems with SSH

- ▶ Byte stream is wrong layer of abstraction for screen.
 - ▶ Want client at *latest* screen.
 - ▶ SSH doesn't understand data, so must send everything.
- ▶ Typing and editing on high-latency path is frustrating.


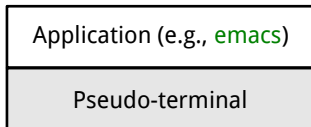
Solution 1: State Synchronization Protocol

- ▶ Runs over UDP.
- ▶ Instead of synchronizing *octet streams*, synchronize *objects*.
- ▶ Object represents state of endpoint.
- ▶ Implements simple interface:
 - ▶ diff: make vector from state $A \rightarrow B$
 - ▶ patch: apply vector to A , producing B


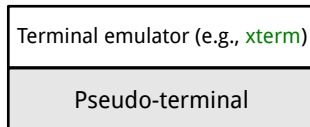
State Synchronization Protocol (cont.)

- ▶ Protected by AES-OCB (Krovetz 2011)
- ▶ Key exchange happens out of band.
 - ▶ Uses SSH to bootstrap.
 - ▶ No privileged code, no daemons.
- ▶ Roaming is easy:
 - ▶ Source address of latest authentic datagram from client \Rightarrow new destination address for server.


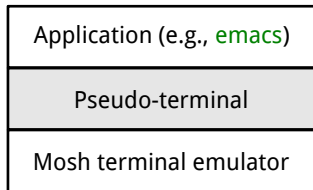
SSH Server



SSH Client





Mosh Server



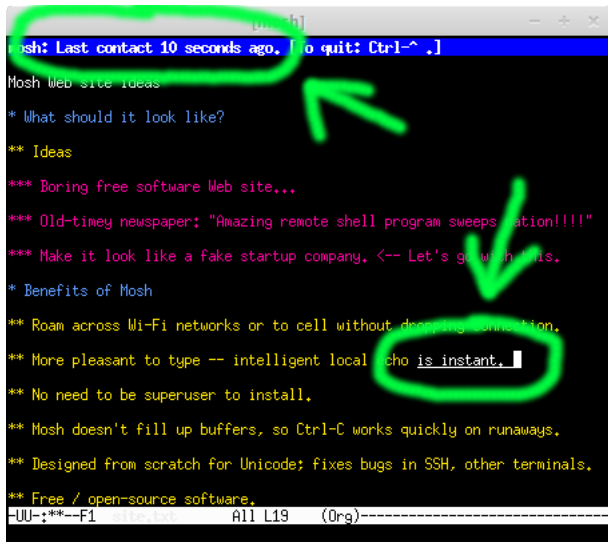
Synced
objects

Mosh Client



Solution 2: Speculative Local Echo and Editing

- ▶ Client anticipates server response.
- ▶ Runs predictive model in the background.
 - ▶ If user hits keystroke, predict key will appear where cursor was.
- ▶ Make predictions in *epochs*.
- ▶ If any prediction from epoch n is confirmed, show all predictions made in that epoch.
- ▶ If user does something difficult to handle, become *tentative* by incrementing epoch.



The screenshot shows a terminal window titled "[mosh]". The first line is a status bar: "mosh: Last contact 10 seconds ago. [to quit: Ctrl-^ .]". This line is highlighted in blue. Below it, the text "Mosh Web site Ideas" is displayed. Then, a prompt "*" is followed by "What should it look like?". This is followed by "** Ideas". Then, three lines of "**" status messages: "Boring free software Web site...", "Old-timey newspaper: 'Amazing remote shell program sweeps nation!!!!'", and "Make it look like a fake startup company. <-- Let's go with this.". Then, a prompt "*" is followed by "Benefits of Mosh". Then, four lines of "**" status messages: "Roam across Wi-Fi networks or to cell without dropping connection.", "More pleasant to type -- intelligent local echo is instant.", "No need to be superuser to install.", "Mosh doesn't fill up buffers, so Ctrl-C works quickly on runaways.", "Designed from scratch for Unicode; fixes bugs in SSH, other terminals.", and "Free / open-source software.". The last line is a status bar: "-UU-:***--F1 site.txt All L19 (Org)-----". There are two green annotations: a circle around the first status bar and an arrow pointing to the second status bar, and another circle around the text "echo is instant." with an arrow pointing to it.

```
[mosh]
mosh: Last contact 10 seconds ago. [to quit: Ctrl-^ .]
Mosh Web site Ideas
* What should it look like?
** Ideas
*** Boring free software Web site...
*** Old-timey newspaper: "Amazing remote shell program sweeps nation!!!!"
*** Make it look like a fake startup company. <-- Let's go with this.
* Benefits of Mosh
** Roam across Wi-Fi networks or to cell without dropping connection.
** More pleasant to type -- intelligent local echo is instant.
** No need to be superuser to install.
** Mosh doesn't fill up buffers, so Ctrl-C works quickly on runaways.
** Designed from scratch for Unicode; fixes bugs in SSH, other terminals.
** Free / open-source software.
-UU-:***--F1 site.txt All L19 (Org)-----
```

Demo


Benefits

- ▶ Roaming and suspend/resume:
 - ▶ Sleep and wake up later.
 - ▶ Change networks at will (Wi-Fi, cellular, wired, VPN).
- ▶ Helpful warnings:
 - ▶ ... if displayed state is stale.
 - ▶ ... if downlink working but not uplink.
- ▶ Interactive over flaky paths.
 - ▶ Works even across 50% loss paths.
 - ▶ Good interactivity even when RTT is > 100 ms.
 - ▶ Semantically appropriate flow control (won't fill up queues, Ctrl-C always works, no beeping fits).
- ▶ Security
 - ▶ Uses SSH to bootstrap: no privileged code, no daemons.
 - ▶ AES-OCB
- ▶ Better Unicode support.

Evaluation

- ▶ Collected 40 hours of terminal usage from six users.
- ▶ Covers 10,000 keystrokes using shell, e-mail, text editor (emacs and vi), chat, Web browser.
- ▶ Replayed over:
 - ▶ Sprint 1xEV-DO (3G)
 - ▶ Verizon LTE (4G)
 - ▶ MIT-Singapore
 - ▶ 50% loss path
- ▶ Result: 70% of keystrokes predicted instantly.
- ▶ Prediction errors $< 1\%$

Sprint 1xEV-DO cumulative keystroke response distribution



Evaluation (cont.)

Verizon LTE service in Cambridge, Mass., running one concurrent TCP download:

	Median latency	Mean	σ
SSH	5.36 s	5.03 s	2.14 s
Mosh	< 0.005 s	1.70 s	2.60 s

Deployment


- ▶ Distributed in Debian, Ubuntu, Fedora, Gentoo, Arch, Slackware versions of GNU/Linux.
- ▶ Available via EPEL for Red Hat, CentOS, Oracle Linux.
- ▶ Included in MacPorts, Homebrew, FreeBSD ports collections.
- ▶ Works on Cygwin and Solaris, (very raw) on Android and iOS.
- ▶ News stories in April on Hacker News, Reddit, The Register, Twitter, Slashdot, Barrapunto.
- ▶ Top repository of the month on GitHub.
- ▶ 200,000+ page views, 50,000+ downloads, 1,200 “followers” of version control repo.

Reception

- ▶ “one of those times you don’t realize something is broken until you see it fixed” — [@xlfe](#)
- ▶ “If you are an SSH user, check out mosh.mit.edu - the user experience really is dreamy.” — [@adamhjk](#)
- ▶ “mosh is awesome. Tested it for two weeks and it really made my life easier: faster feedback and no more reconnects(!)” — [@esmolanka](#)
- ▶ “Finalement, la vie d’admin c’est pas si Mosh que ça” — [@korben](#)
- ▶ “There is very (if any) little research content.” — [USENIX review](#)
- ▶ “ISO 2022 locking escape sequences oh flying spaghetti monster please kill me now.” — [USENIX review](#)

State Sync Protocol for all?

- ▶ We believe SSP may be appropriate for many network problems.
- ▶ Android Gmail, Google Chat cannot roam without failure.
- ▶ **May 15, 2012:**



- ▶ Neither can Gmail (Web edition).
- ▶ These problems can be expressed as state synchronization.

Next Steps

- ▶ Mosh paper to be presented at USENIX ATC (June 2012).
- ▶ Essay to appear in *;login:* magazine.
- ▶ Mosh software under development by a team of contributors.
- ▶ We are working to apply SSP to mobile videoconferencing.
- ▶ We hope to show quantitative improvement on standard metrics (latency, quality), plus features like roaming.

Summary

- ▶ SSP is a secure datagram protocol that synchronizes abstract objects across a roaming IP connection.
- ▶ Mosh uses SSP to synchronize a terminal emulator with predictive local echo.
- ▶ In evaluations with 10,000 real-world keystrokes from six users, Mosh markedly reduced user-visible latency across several Internet paths.
- ▶ We think SSP will be useful for other applications as well.
- ▶ <http://mosh.mit.edu>