

# Mosh

## An Interactive Remote Shell for Mobile Clients

Keith Winstein (with Hari Balakrishnan)

[keithw@mit.edu](mailto:keithw@mit.edu)

May 16, 2012

# Secure Shell, 1996

- ▶ Connects local terminal to remote terminal.
- ▶ Conveys over TCP:
  - ▶ user keystrokes  $\Rightarrow$  server
  - ▶ octet stream (coded screen updates)  $\Rightarrow$  client terminal
- ▶ Connection endpoints dictated by IP:port on both sides

# Post-1996 problems with SSH

- ▶ Can't roam:
  - ▶ ... across Wi-Fi networks.
  - ▶ ... from Wi-Fi to cell or vice versa.
- ▶ TCP times out if data unacknowledged after  $n$  minutes.
  - ▶ Session dies if laptop goes to sleep.
- ▶ TCP responds poorly to packet loss.

# More problems with SSH

- ▶ Byte stream is wrong layer of abstraction for screen.
  - ▶ If client screen state is old, want to skip ahead directly.
  - ▶ Don't want to replay megabytes in between.
  - ▶ But SSH doesn't know how client interprets octets, so must send all.
  - ▶ TCP fills buffers, so bufferbloat means Control-C takes forever.
- ▶ Typing and editing on high-latency path is frustrating.
  - ▶ Cellular wireless (100 ms to 500 ms)
  - ▶ Intercontinental (250 ms)
  - ▶ Loaded 4G LTE (5,000 ms!)

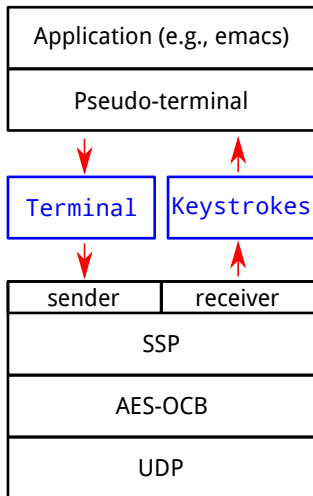
# Solution 1: State Synchronization Protocol

- ▶ Runs over UDP.
- ▶ Instead of synchronizing *octet streams*, synchronize *objects*.
- ▶ Object represents state of endpoint.
- ▶ Implements simple interface:
  - ▶ diff: make vector from state  $A \rightarrow B$
  - ▶ patch: apply vector to  $A$ , producing  $B$
  - ▶ subtract: remove common prefix from  $A, B$
- ▶ Object implementation, **not protocol**, defines synchronization semantics.

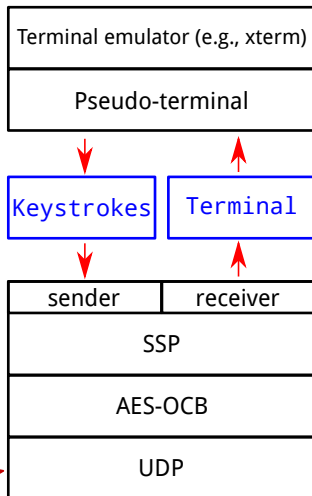
# State Synchronization Protocol (cont.)

- ▶ Runs over UDP
- ▶ Protected by AES-OCB (Krovetz 2011)
  - ▶ Integrity and confidentiality with one key.
- ▶ Roaming is easy:
  - ▶ Source address of latest authentic datagram from client  $\Rightarrow$  new destination address for server.

## Mosh Server



## Mosh Client



# State Synchronization Protocol (cont.)

- ▶ **Flow control:** try to have at-most one diff in flight.
- ▶ Use TCP's SRTT/retransmission timeout with tweaks.
- ▶ Minimum interval between frames is  $SRTT/2$ .
- ▶ Can skip intermediate states.
- ▶ “Pretransmissions” (post-paper improvement).



# Pretransmissions (post-paper improvement)

“Prophylactic” retransmission reduces latency in presence of loss.

1. Last ack was for state #3. Then state changes to #4.
2. Host sends diff from  $3 \rightarrow 4$ .
3. Object state changes to state #5.
4. If last diff hasn't timed out, formulate next diff as  $4 \rightarrow 5$ .
5. **Also** make diff from  $3 \rightarrow 5$ : the *pretransmission*.
6. If retransmission is shorter or not “much” longer, send it instead.

## Solution 2: Different semantics for different directions

- ▶ Client  $\Rightarrow$  server: object represents *history of keystrokes typed*.
  - ▶ Same semantics as TCP.
  - ▶ `subtract` ensures only outstanding data actually stored in memory.
- ▶ Server  $\Rightarrow$  client: object represents *current screen state*.
  - ▶ *E.g.* 80x24 terminal, Unicode grapheme in each cell + bell.
  - ▶ Okay to skip over intermediate states!
  - ▶ Server and client **both have model** of what's on the screen.

# Benefits

- ▶ Can sleep and wake up.
- ▶ Can roam across ESSIDs or to cell network.
- ▶ Can fly to Canada and open laptop in hotel.
- ▶ Can warn user when displayed state is stale.
- ▶ Semantically appropriate flow control won't fill up network queues.
  - ▶ Control-C always works within RTT.
  - ▶ No beeping fits (one beep per diff).
- ▶ Supports lossy links.
- ▶ Uses SSH to start connection.
  - ▶ No privileged (root) code!
  - ▶ No daemons.
- ▶ Better Unicode support.

# Unicode admits varying interpretations.

```
xterm 271
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
hello
sh$
```

```
[mosh]
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
hello
sh$
```

```
GNOME Terminal 3.0.1
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
hello
sh$
```

```
Macintosh HD — Terminal.app 2.2.2
sh$ echo -e "xyz\033[2;2H\0314\0202\nhello"
xyz
ello
sh$
```

**bricks the terminal!**

## Solution 3: Speculative Local Echo and Editing


- ▶ Now that client has complete picture of screen state, can start to anticipate server response.
- ▶ Runs predictive model in the background.
  - ▶ If user hits keystroke, predict key will appear where cursor was.
- ▶ Make predictions in *epochs*.
- ▶ If any prediction from epoch  $n$  is confirmed, show all predictions made in that epoch.
  - ▶  $n$  might be current epoch or one from the past.
- ▶ If user does something difficult to handle, become *tentative* by incrementing epoch.
  - ▶ Moves to new row.
  - ▶ Hits control character (including carriage return).
  - ▶ Uses up/down arrow key.

# Demo

# Evaluation

- ▶ Collected 40 hours of terminal usage from six anonymous users.
- ▶ Covers 9,986 keystrokes using shell, e-mail, text editor (emacs and vi), chat, Web browser.
- ▶ Replayed sessions over Sprint 1xEV-DO (3G) commercial network over SSH and Mosh.
- ▶ Result: 70% of keystrokes can be predicted instantly.
- ▶ Prediction errors  $< 1\%$ 
  - ▶ Most common cause: word wrap
  - ▶ Bad prediction removed from display within RTT.
- ▶ Median UI latency **503 ms** (SSH) to  **$< 5$  ms** (Mosh).
- ▶ Mean UI latency **515 ms** (SSH) to **173 ms** (Mosh).

# Keystroke response times, cumulative distribution





## Evaluation (cont.)

**Verizon LTE service in Cambridge, Mass., running one concurrent TCP download:**

	Median latency	Mean	$\sigma$
SSH	5.36 s	5.03 s	2.14 s
Mosh	< 0.005 s	1.70 s	2.60 s

**MIT-Singapore Internet path (to Amazon EC2 data center):**

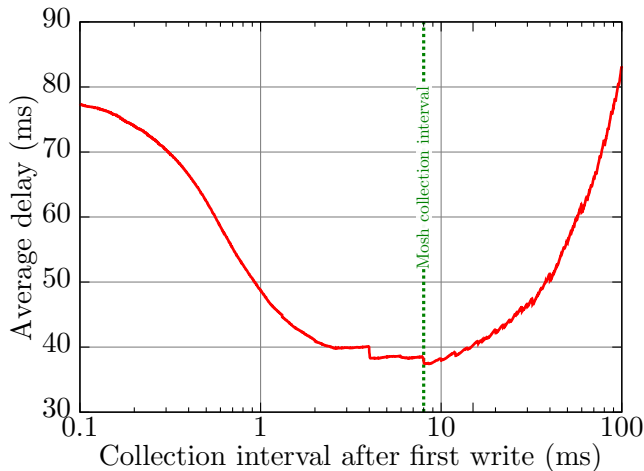
	Median latency	Mean	$\sigma$
SSH	273 ms	272 ms	9 ms
Mosh	< 5 ms	86 ms	132 ms

# SSP with high packet loss

**Synthetic link with 100 ms RTT, 50% round-trip  
i.i.d. packet loss:**

	Median	Mean	$\sigma$
SSH	0.416 s	16.8 s	52.2 s
Mosh (no predictions)	0.222 s	0.329 s	1.63 s

## Average protocol-induced delay from varying collection interval (SRTT > 500 ms)



# Deployment


- ▶ Distributed in Debian, Ubuntu, Fedora versions of GNU/Linux.
- ▶ Included in MacPorts, Homebrew, FreeBSD ports collections.
- ▶ Web site at `mosh.mit.edu`
- ▶ News stories in April on Hacker News, Reddit, The Register, Twitter, Slashdot, Barrapunto.
- ▶ 130,000 unique visitors, 25,000+ downloads, 1,100+ “followers” of source code repository on Github.

# Reception

- ▶ “one of those times you don’t realize something is broken until you see it fixed” — [@xlfe](#)
- ▶ “If you are an SSH user, check out mosh.mit.edu - the user experience really is dreamy.” — [@adamhjk](#)
- ▶ “mosh is awesome. Tested it for two weeks and it really made my life easier: faster feedback and no more reconnects(!)” — [@esmolanka](#)
- ▶ “Finalement, la vie d’admin c’est pas si Mosh que ça” — [@korben](#)
- ▶ “There is very (if any) little research content.” — [USENIX review](#)
- ▶ “ISO 2022 locking escape sequences oh flying spaghetti monster please kill me now.” — [USENIX review](#)

# State Sync Protocol for all?

- ▶ We believe SSP may be appropriate for many network problems.
- ▶ Android Gmail, Google Chat cannot roam without failure.
- ▶ **May 15, 2012:**



- ▶ Neither can Gmail (Web edition).
- ▶ These problems can be expressed as state synchronization.

# Next Steps

- ▶ Mosh paper to be presented at USENIX ATC (June 2012).
- ▶ Mosh software under development by a team of contributors.
- ▶ We are working to apply SSP to mobile videoconferencing.
- ▶ We hope to show quantitative improvement on standard metrics (latency, quality), plus features like roaming.

# Summary

- ▶ SSP is a secure datagram protocol that synchronizes abstract objects across a roaming IP connection.
- ▶ Mosh uses SSP to synchronize a terminal emulator with predictive local echo.
- ▶ In evaluations with 10,000 real-world keystrokes from six users, Mosh markedly reduced user-visible latency across several Internet paths.
- ▶ We think SSP will be useful for other applications as well.
- ▶ <http://mosh.mit.edu>